

SPR Lab 3

March 3, 2021

1 Multivariate Bernoulli Distribution

Run EM and K-Means algorithm on same dataset and compare the results. *## Reading MNIST dataset*

```
[1]: import numpy as np
import torch
import pytorch_lightning as pl
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import scipy
import seaborn as sns
np.random.seed(42)

class Config:
    N = 600
    K = 3
    D = 784
    iterations = 50

def load_mnist():
    mnist = pd.read_csv('mnist_train.csv').values
    X_train, y_train = mnist[:, 1:], mnist[:, 0]

    # take only digits 2, 3 and 4
    take = (y_train >= 2) & (y_train <= 4)
    X_train, y_train = X_train[take], y_train[take]

    # making arrays have only zeros or ones
    X_train = ((X_train/255.0) > .5).astype(int)

    # make sure you have a total of 600 examples
    train_size = Config.N / y_train.shape[0]
    X_train, _, y_train, _ = train_test_split(X_train, y_train,
    ↪train_size=train_size, stratify=y_train, random_state=42)


    # show one example from each class
    fig, axs = plt.subplots(1, 3, figsize=(18, 6))
```

```

for i, c in enumerate([2, 3, 4]):
    img = X_train[y_train == c][0].reshape((28, 28))
    t = (y_train==c).sum()
    axs[i].imshow(img, cmap='CMRmap')
    axs[i].set(xticks = [])
    axs[i].set_xlabel(f"Class={c}, {t} examples", fontsize=14)
plt.show()
return X_train, y_train

```

```
X_train, y_train = load_mnist()
```



1.1 EM algorithm for Multivariate Bernoulli Mixture

One iteration of EM updates π_k and μ_k as follows

$$P(x|\mu_k) = \prod_{i=1}^D \mu_{ki}^{x_i} (1 - \mu_{ki})^{(1-x_i)}$$

$$\gamma_{nk} = \frac{\pi_k P(x_n|\mu_k)}{\sum_{j=1}^K \pi_j P(x_n|\mu_j)}$$

$$N_k = \sum_{n=1}^N \gamma_{nk}$$

$$\mu_k = \frac{\sum_{n=1}^N \gamma_{nk} x_n}{N_k}$$

```

[2]: class BernoulliMixtureModel:
    def __init__(self):
        self.pi = np.ones((Config.K,1)) / Config.K # pi_k
        self.mu = np.random.uniform(low=.25, high=.75, size=(Config.K, Config.
        ↪D)) # mu_kj
        self.mu /= self.mu.sum(axis=1, keepdims=True) # sum mu_kj over j should
        ↪be 1

    # helpers
    def prob(self, X_nd, k): # examples, k -> [P(example | k)]
        mu_1d = self.mu[k:k+1]
        p_nd = mu_1d**X_nd * (1- mu_1d)**(1-X_nd) #(1, d)*(n, d)

```

```

p_n = np.prod(p_nd, axis=1) # product over dimensions, axis=1
return p_n

def getGamma(self, X_nd):
    g_kn = np.array([
        self.pi[k, 0] * self.probab(X_nd, k)
        for k in range(Config.K)
    ]) # (k n)

    denom = g_kn.sum(axis=0, keepdims=True)
    denom[denom == 0.0] = 1
    g_kn /= denom # sum over clusters (k) must be one

    return g_kn, denom

def oneIteration(self, X_nd):
    """Performs one iteration of the EM algorithm for bernoulli mix model
    Args
        * X_nd is an np.array with n examples, each of d dimensions
    Function
        * updates Params.mu and Params.pi
    """
    g_kn, denom = self.getGamma(X_nd)

    # N_k is gamma summed over n
    N_k = g_kn.sum(axis=1, keepdims=True) # shape (k,)

    # new mu is examples averaged over gamma
    mu_kd = np.matmul(g_kn, X_nd)
    self.mu = mu_kd / N_k
    self.pi = N_k / Config.N

    # return log likelihood
    return - np.log(denom).sum()

def fit(self, X_train):
    return [
        self.oneIteration(X_train)
        for i in range(Config.iterations)
    ]


```

```

[3]: model = BernoulliMixtureModel()
neg_log_likelihoods = model.fit(X_train)
fig, ax = plt.subplots(figsize=(10, 8))
ax.plot(range(10, Config.iterations), neg_log_likelihoods[10:])
ax.set_ylabel('log likelihood', fontsize=16)
ax.set_xlabel('iteration', fontsize=16)

```

```
ax.set_title('EM algorithm minimizing negative log likelihood', fontsize=20);
```



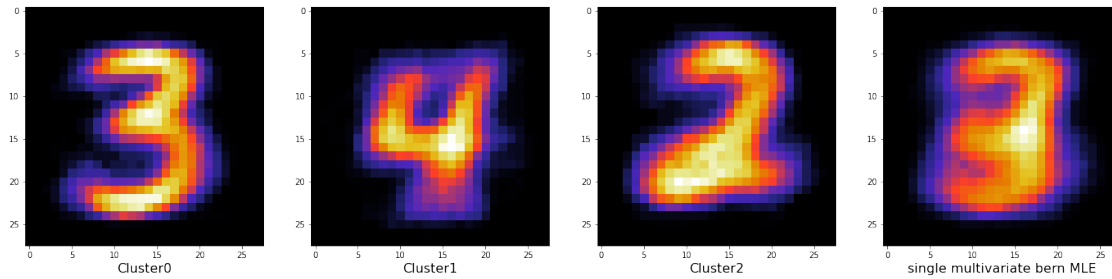
1.2 Most likely estimate for Single Multivariate Bernoulli

$$Z = [z_i] \in R^d$$

$$Z_{mle} = \frac{\sum_{n=1}^N X_i}{N}$$

```
[4]: # display EM algo mus
images = (model.mu * 255).reshape((Config.K, 28, 28)).astype(np.uint8)
fig, axs = plt.subplots(1, 4, figsize=(24, 6))
for k in range(Config.K):
    axs[k].imshow(images[k], cmap='CMRmap')
    axs[k].set_xlabel(f'Cluster{k}', fontsize=16)

img = (X_train.mean(axis=0) * 255).reshape((28, 28)).astype(np.uint8)
axs[Config.K].imshow(img, cmap='CMRmap')
axs[Config.K].set_xlabel('single multivariate bern MLE', fontsize=16)
plt.show()
```



1.3 K-means clustering for MNIST

```
[5]: class Kmeans:
    # standardize data
    def standardize(self, X_train):
        sigma = X_train.std(axis=0, keepdims=True)
        sigma[sigma == 0] = 1
        X_stan = X_train / sigma
        return X_stan, sigma

    def fit(self, X_train):
        # init means as K random points from data
        X_stan, sigma = self.standardize(X_train)
        idx = np.random.choice(np.arange(Config.N), Config.K)
        self.mu = X_stan[idx] * 1.0

        # init assign
        z_n = np.random.choice(np.arange(Config.K), Config.N)

        losses = []
        for i in range(Config.iterations):
            # calc pairwise dists (ND - K1D -> KN shapes for ref)
            d_kn = np.linalg.norm(X_stan - self.mu[:,None,:], axis=-1)
            # assign cluster based on proximity
            z_n = d_kn.argmin(axis=0) # N
            losses.append(d_kn.min(axis=0).sum())

            # update cluster means
            z_kn = np.eye(3)[: , z_n] # 3 N
            z_kn /= z_kn.sum(axis=1, keepdims=True) # normalize
            self.mu = z_kn @ X_stan

        self.mu = self.mu * sigma
        return losses

    def predict(self, X):
```

```

d_kn = np.linalg.norm(X - self.mu[:,None,:], axis=-1)
# assign cluster based on proximity
z_n = d_kn.argmin(axis=0) # N
return z_n

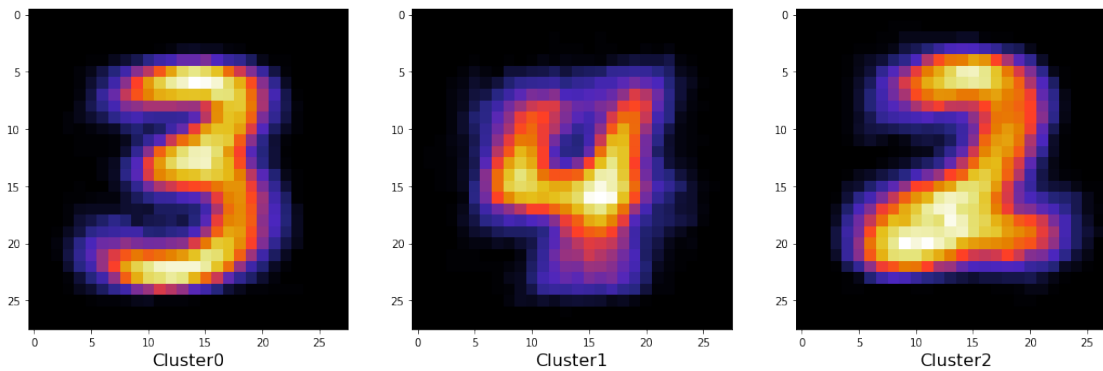
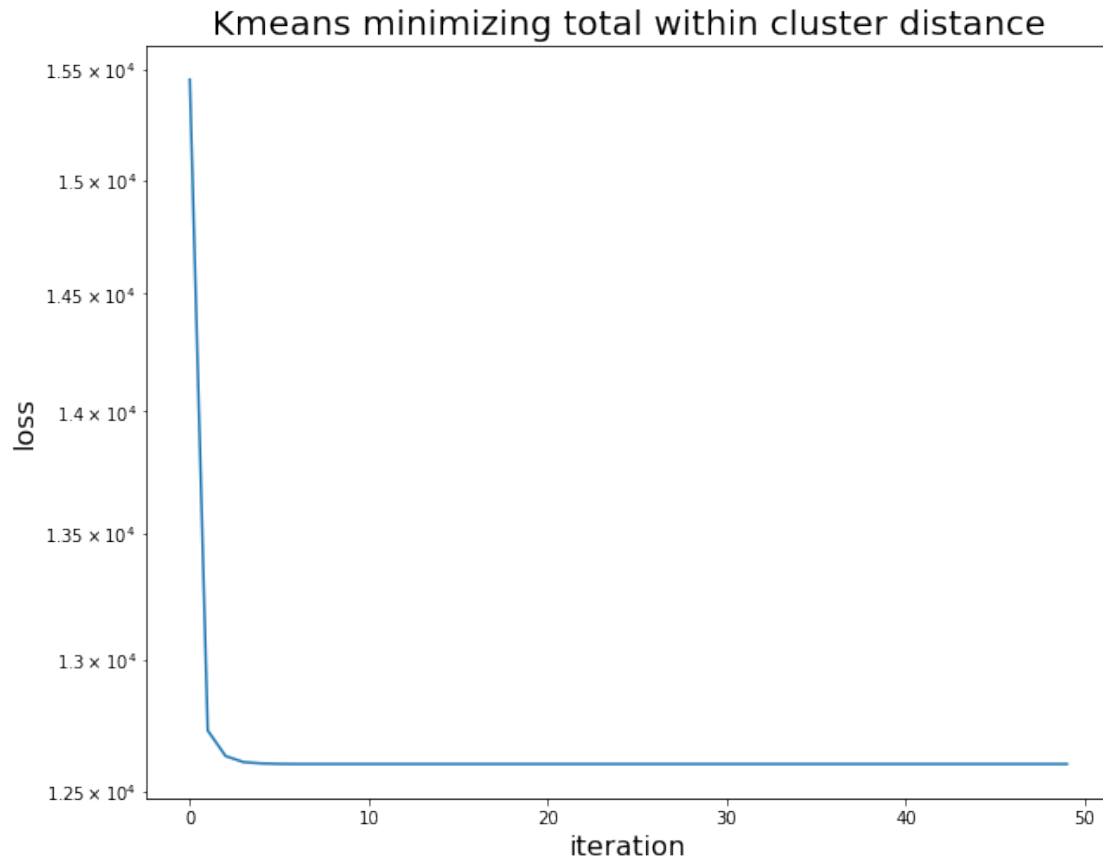
def plot(self, losses, display_means=True):
    fig, ax = plt.subplots(figsize=(10, 8))
    ax.plot(range(Config.iterations), losses)
    ax.set_ylabel('loss', fontsize=16)
    ax.set_xlabel('iteration', fontsize=16)
    ax.set_title('Kmeans minimizing total within cluster distance',
→fontsize=20)
    plt.yscale('log')
    plt.show()

    # display the k means
    if display_means:
        images = (model.mu * 255).reshape((Config.K, 28, 28)).astype(np.
→uint8)

        fig, axs = plt.subplots(1, 3, figsize=(18, 6))
        for k in range(Config.K):
            axs[k].imshow(images[k], cmap='CMRmap')
            axs[k].set_xlabel(f'Cluster{k}', fontsize=16)
        plt.show()

model = Kmeans()
losses = model.fit(X_train)
model.plot(losses)

```



2 GMM on custom data

Sample data from univariate and multivariate Gaussian Mixture Models and verify whether EM algorithm is able to retrieve the mixture densities in each of the case.

```
[6]: # Config has hardcoded values and are not meant to be changed at all for this_
      ↪notebook

class Config:
    N = 1700
    K = 3 # theser
    D = 2 # dont change this we need 2d for interpreting plots
    iterations = 75
    eps = 1e-3

class Actual:
    mu = np.array([[3,3], [10, -10], [5, 20]])
    cov_norot = [
        [[4, 0], [0, 64]],
        [[16, 0], [0, 1]],
        [[16, 0], [0, 6]]
    ]
    cov = []
    rots = [45, 10, 30]
    sizes = [500, 200, 1000]

def rotate2D(theta):
    c, s = np.cos(theta), np.sin(theta)
    return np.array([[c, -s], [s, c]])

def sampleMultiGaussian(mu, cov, theta, size):
    R = rotate2D(theta)
    Actual.cov.append(R.T @ cov @ R)
    return np.random.multivariate_normal(mu, Actual.cov[-1], size=size)

def sampleDistriubtion():
    fig, ax = plt.subplots(figsize=(10, 10))
    X_train = []
    y_train = []

    x_grid = np.arange(-20, 25, .02)
    y_grid = np.arange(-20, 27, .02)
    x_grid, y_grid = np.meshgrid(x_grid, y_grid)
    coordinates = np.array([x_grid.ravel(), y_grid.ravel()]).T
    colors = ['blue', 'orange', 'green']

    for i in range(Config.K):
        samples = sampleMultiGaussian(Actual.mu[i], Actual.cov_norot[i],
                                      theta=np.radians(Actual.rots[i]), size=Actual.sizes[i])
        labels = np.ones(Actual.sizes[i]) * i
        X_train.append(samples)
        y_train.append(labels)
        ax.scatter(samples[:,0], samples[:, 1], label=str(i+1))
```



```

        z_grid = scipy.stats.multivariate_normal(Actual.mu[i], Actual.cov[i]).
→pdf(coordinates).reshape(x_grid.shape)
        ax.contour(x_grid, y_grid, z_grid, colors=colors[i])

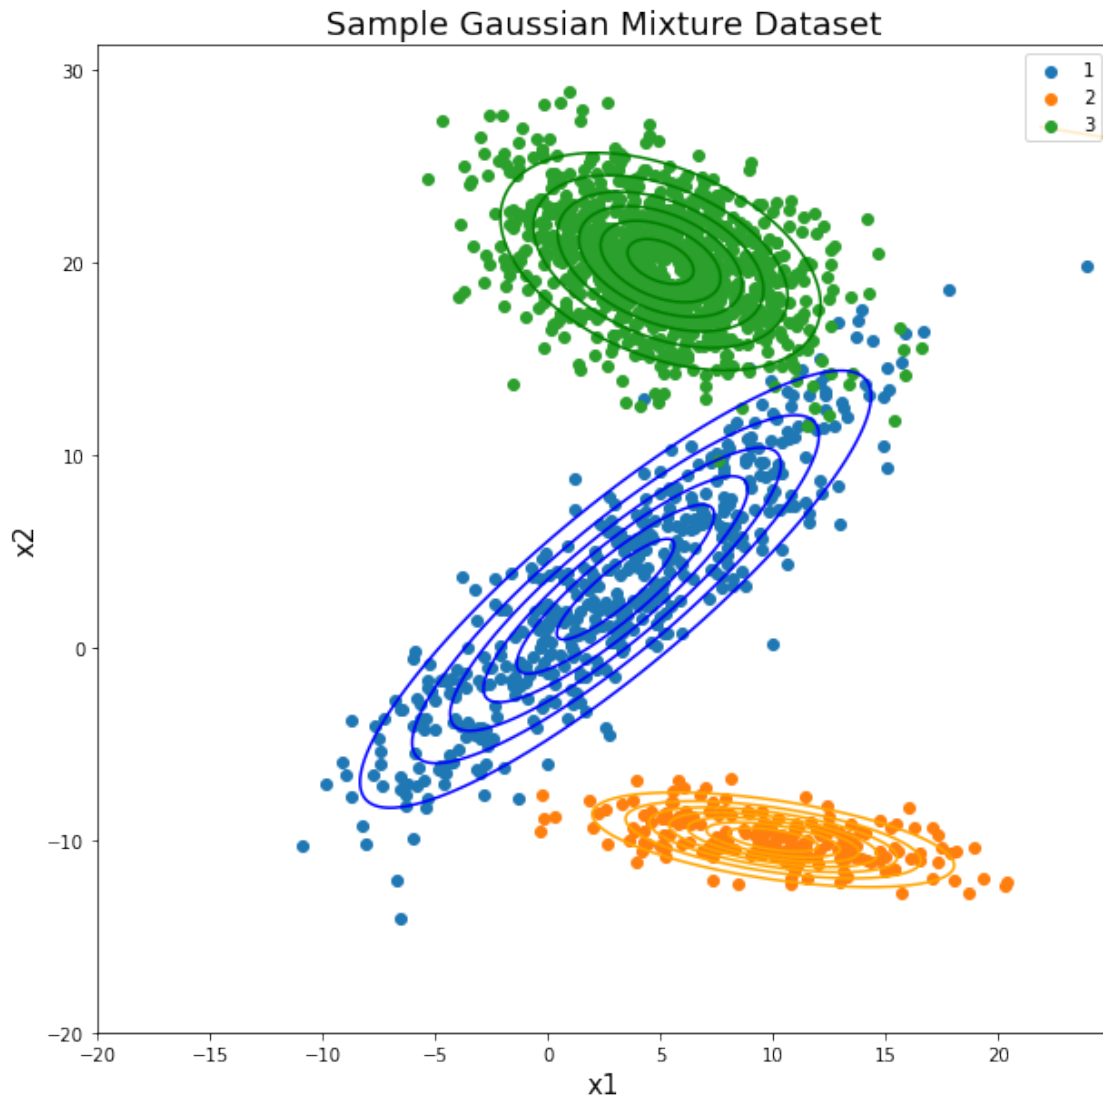
    ax.set_xlabel('x1', fontsize=15)
    ax.set_ylabel('x2', fontsize=15)

    ax.set_title('Sample Gaussian Mixture Dataset', fontsize=18)
    ax.legend()

    return np.concatenate(X_train, axis=0), np.concatenate(y_train, axis=0)

X_train, y_train = sampleDistriubtion()

```



```
[7]: class GaussianMixtureModel:
    def fit(self, X_train):
        self.pi = np.ones((Config.K,1)) / Config.K
        ↪ # pi_k
        idx = np.random.choice(np.arange(Config.N), Config.K)
        self.mu = X_train[idx] * 1.0 # choose K samples randomly from dist
        self.cov = np.array([np.eye(Config.D) for i in range(Config.K)]) # init
        ↪ K covs to identity

        nll = []
        for i in range(Config.iterations):
            loss = self.em_one_iteration(X_train)
            nll.append(loss)
```

```

        if len(nll) > 1 and np.abs(loss - nll[-2]) < Config.eps:
            break
    return nll

# helpers
def prob(self, X_nd, k): # (examples, k) -> [P(example1 | k),
    ↪ P(example2|k), ..]
    mu_1d = self.mu[k:k+1]
    cov_dd = self.cov[k]
    z_nd = X_nd - mu_1d
    A_nd = (z_nd @ np.linalg.inv(cov_dd))
    z_n = (A_nd * z_nd).sum(axis=1)
    p_n = (2*np.pi) ** (- Config.D / 2) * np.linalg.det(cov_dd) ** -.5 *
    ↪ np.exp(-.5 * z_n)
    return p_n

def getGamma(self, X_nd):
    g_kn = np.array([
        self.pi[k, 0] * self.prob(X_nd, k)
        for k in range(Config.K)
    ]) # (k n)
    denom = g_kn.sum(axis=0, keepdims=True)
    denom[denom == 0.0] = 1
    g_kn /= denom # sum over clusters (k) must be one

    return g_kn, denom

def em_one_iteration(self, X_nd):
    """Performs one iteration of the EM algorithm for gaussian mixture model
    Args
        * X_nd is an np.array with n examples, each of d dimensions
    Function
        * updates Params.mu, Params.cov and Params.pi
    """
    g_kn, denom = self.getGamma(X_nd)

    # N_k is gamma summed over n
    N_k1 = g_kn.sum(axis=1, keepdims=True) # shape (k,1)

    # new mu is examples averaged over gamma
    self.mu = np.matmul(g_kn, X_nd) / N_k1
    self.pi = N_k1 / Config.N

    # new mean is used to calculate new cov
    # calculate cov as E[x-mu] but over dist of reponse coeff
    for k in range(Config.K):
        g_n1 = g_kn[k:k+1].T

```

```


        z_nd = X_nd - self.mu[k:k+1]
        self.cov[k] = ((g_n1 * z_nd).T @ z_nd) / N_k1[k]

    # return log likelihood
    neg_log_likelihood = - np.log(denom).sum()
    return neg_log_likelihood

model = GaussianMixtureModel()
nll = model.fit(X_train)
fig, ax = plt.subplots(figsize=(8, 6))
plt.plot(nll)
plt.xlabel('iterations')
plt.ylabel('Negative log likelihood')
plt.title('Loss after each iteration')
plt.show()

print("Model means:")
print(model.mu)
print()
print("Actual means:")
print(Actual.mu)

```



Model means:

```
[[ 5.00015899 19.8556918 ]
 [ 9.76671785 -9.90188685]
 [ 3.14197768  2.97072897]]
```

Actual means:

```
[[ 3  3]
 [10 -10]
 [ 5 20]]
```

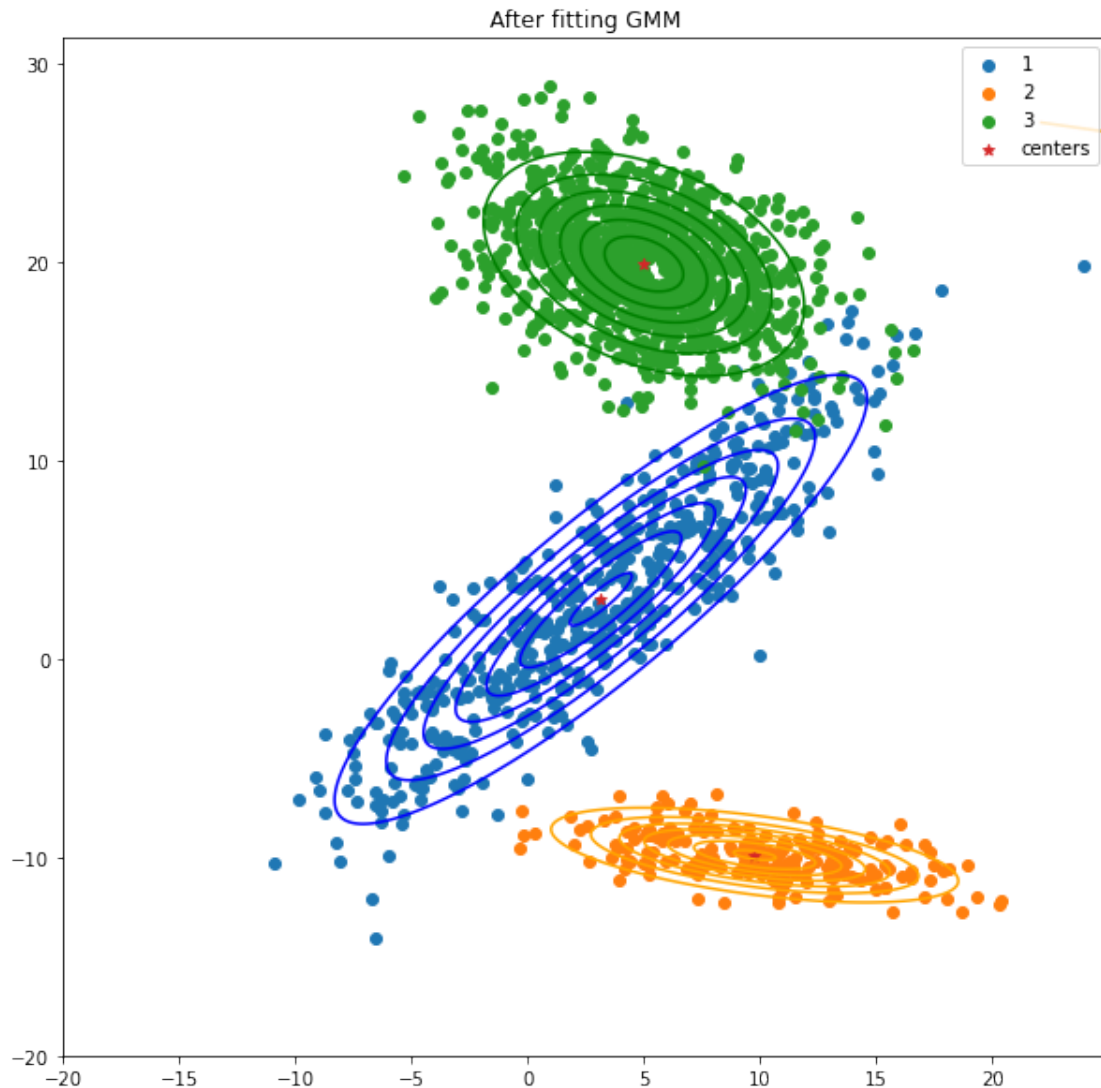
2.1 Plotting learned distributions

```
[8]: fig, ax = plt.subplots(figsize=(10, 10))
parts = np.cumsum([0] + Actual.sizes)

x_grid = np.arange(-20, 25, .02)
y_grid = np.arange(-20, 27, .02)
x_grid, y_grid = np.meshgrid(x_grid, y_grid)
coordinates = np.array([x_grid.ravel(), y_grid.ravel()]).T
colors = ['green', 'orange', 'blue']

for k in range(Config.K):
    x = X_train[parts[k]:parts[k+1]]
    ax.scatter(x[:,0], x[:,1], label=str(k+1))
    z_grid = scipy.stats.multivariate_normal(model.mu[k], model.cov[k]).
    pdf(coordinates).reshape(x_grid.shape)
    ax.contour(x_grid, y_grid, z_grid, colors=colors[k])

ax.scatter(model.mu[:,0], model.mu[:,1], marker='*', label='centers')
ax.set_title('After fitting GMM')
plt.legend();
```



3 Parzen Window and K-means on the same custom data

Run K-Means and Parzen Window/Kernel density estimation on dataset generated in question 2 and compare the results with EM algorithm. Try different window sizes and different kernels for Kernel density estimation. Compare it with histogram based density estimation.

```
[9]: model = Kmeans()
losses = model.fit(X_train)
model.plot(losses, display_means=False)
labels = model.predict(X_train)

# plotting clusters got by model
fig, ax = plt.subplots(figsize=(10, 10))
```

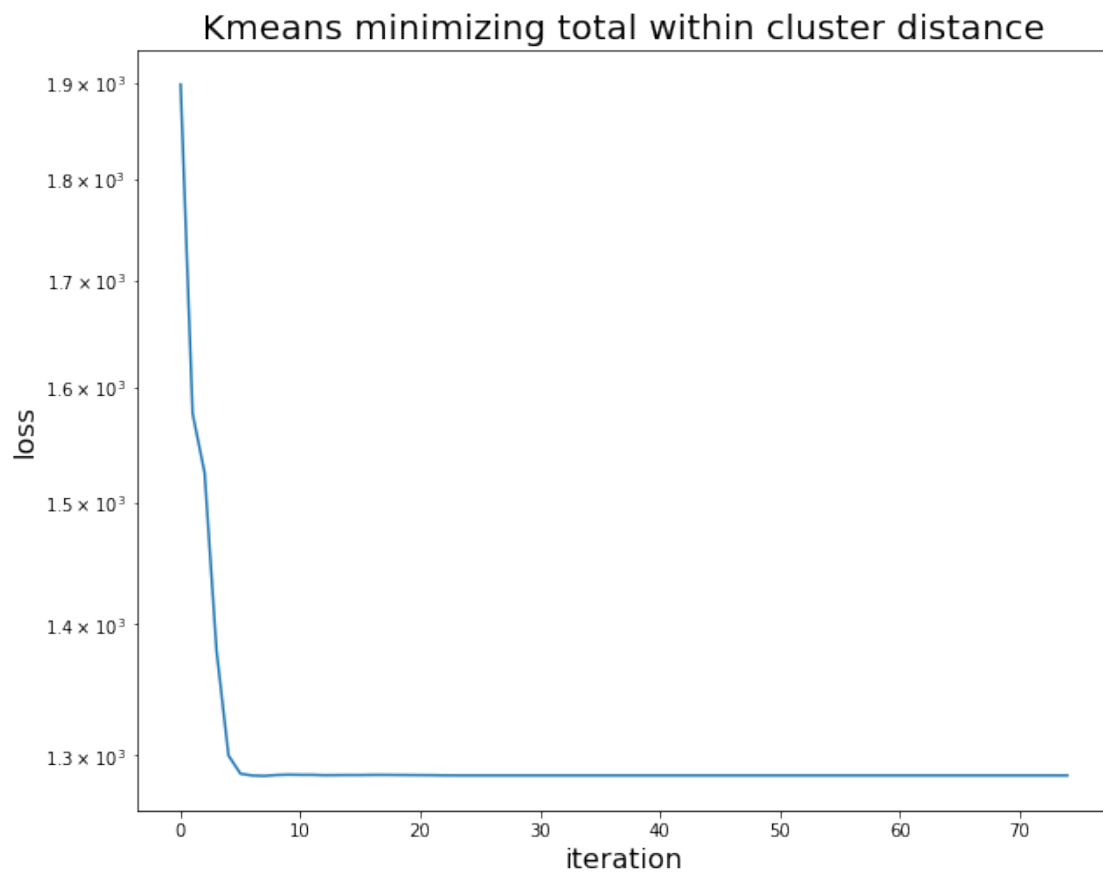
```

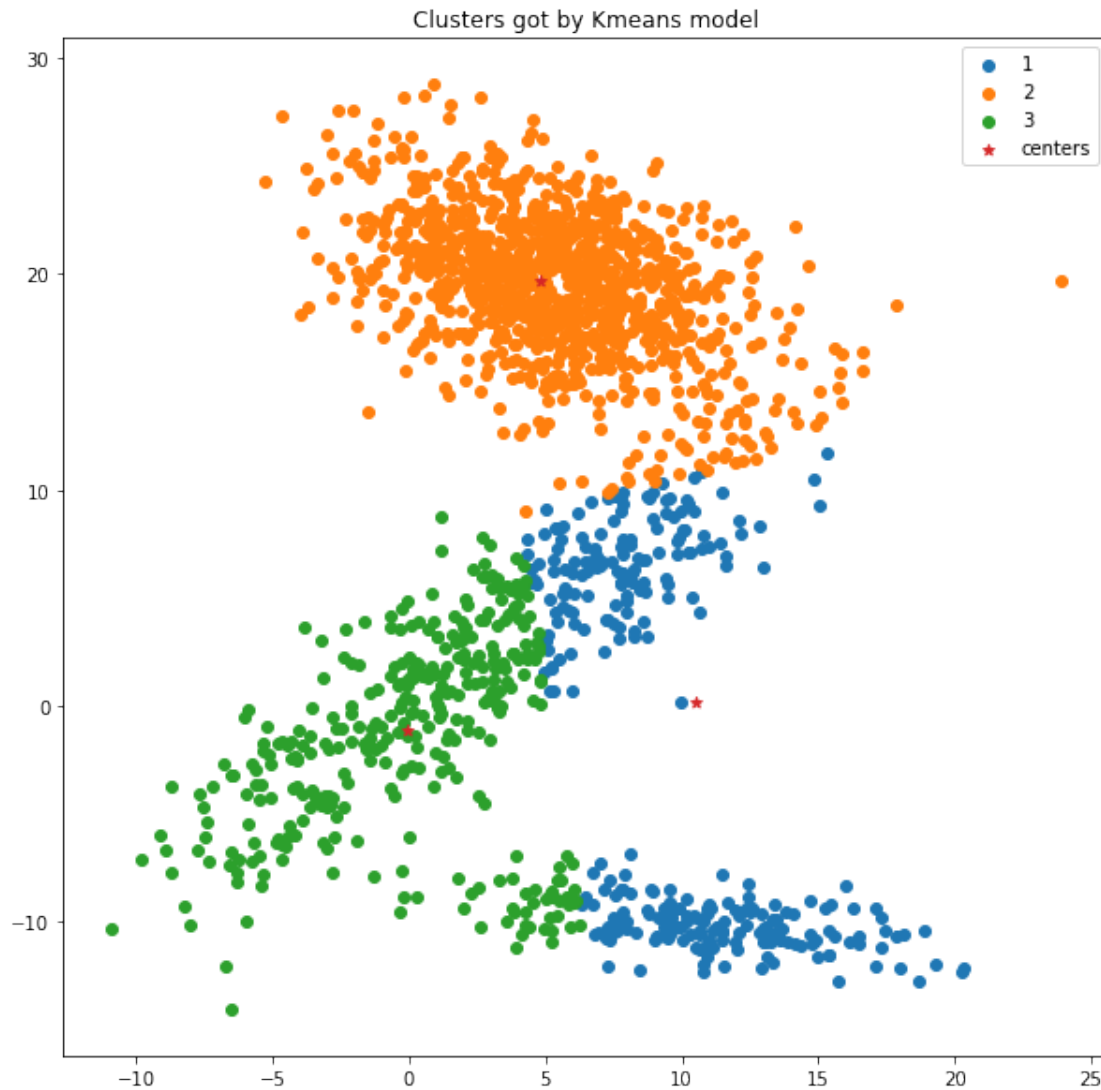
colors = ['blue', 'green', 'orange']

for k in range(Config.K):
    x = X_train[labels == k]
    ax.scatter(x[:,0], x[:,1], label=str(k+1))

ax.scatter(model.mu[:,0], model.mu[:,1], marker='*', label='centers')
ax.set_title('Clusters got by Kmeans model')
plt.legend();

```





```
[10]: def gaussianKernel(z_mtd, h): #  $z = (x - x_i) / h$  of shape  $(m, t, d)$   $m$  train ex,  $t$  for predict arg,  $d = \text{Config.d}$ 
      z_mt = np.linalg.norm(z_mtd, axis=2)
      phi_mt = (2*np.pi) ** (Config.D / 2) * np.exp(- z_mt**2 / 2)
      phi_t = phi_mt.mean(axis=0) / h ** Config.D
      return phi_t

      def windowKernel(z_mtd, h):
          phi_mt = np.prod(z_mtd**2 <= .25, axis=2) # reduce d
          phi_t = phi_mt.mean(axis=0) / h ** Config.D # reduce m
          return phi_t

      class ParzenWindow:
```



```

def __init__(self, X_train, h, kernel='gaussian'):
    self.X = X_train
    self.h = h
    if kernel.lower() == 'gaussian':
        self.kernel = gaussianKernel
    else:
        self.kernel = windowKernel

def predict(self, X_t):
    z_mtd = (X_t - self.X[:,None,:]) / self.h
    phi_t = self.kernel(z_mtd, self.h)
    return phi_t

```

```

[11]: x_grid = np.arange(-20, 25, .5)
      y_grid = np.arange(-20, 33, .5)
      x_grid, y_grid = np.meshgrid(x_grid, y_grid)
      coordinates = np.array([x_grid.ravel(), y_grid.ravel()]).T

      model_gk = ParzenWindow(X_train, 2, 'gaussian')
      density_g = model_gk.predict(coordinates).reshape(x_grid.shape)

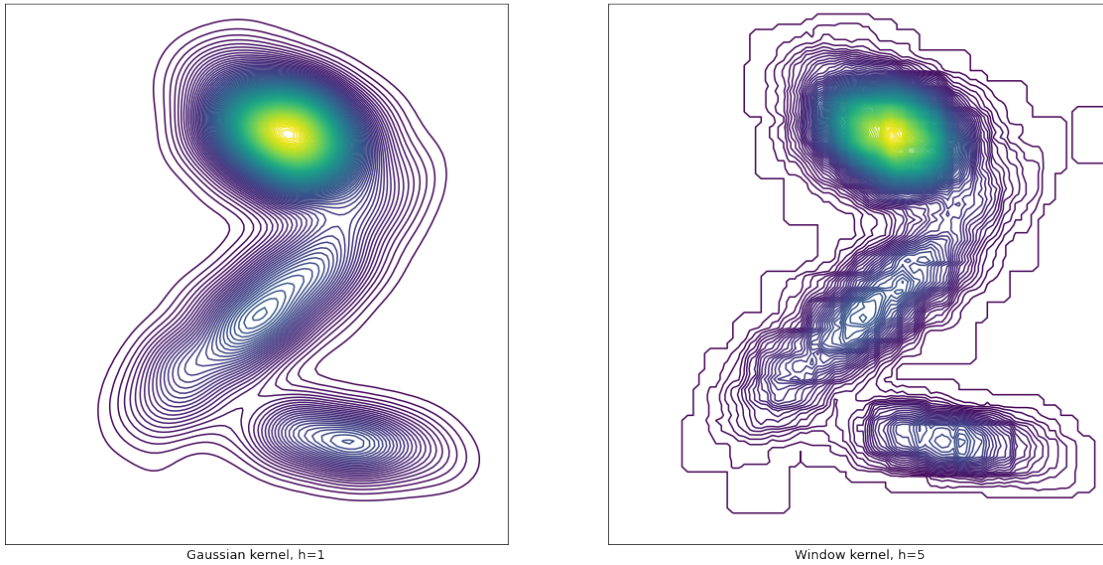
      model_wk = ParzenWindow(X_train, 5, 'window')
      density_w = model_wk.predict(coordinates).reshape(x_grid.shape)

      fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(20, 10))
      fig.suptitle('Parzen window density estimation', fontsize=18)

      ax1.contour(x_grid, y_grid, density_g, levels=np.linspace(density_g.min(),
      ↪density_g.max(), num=100))
      ax1.set_xlabel('Gaussian kernel, h=1', fontsize=13)
      ax1.set(xticks=[], yticks=[])
      ax2.contour(x_grid, y_grid, density_w, levels=np.linspace(density_w.min(),
      ↪density_w.max(), num=100))
      ax2.set_xlabel('Window kernel, h=5', fontsize=13)
      ax2.set(xticks=[], yticks=[]);

```

Parzen window density estimation



```
[12]: h_list = [0.3, 1, 3, 9]

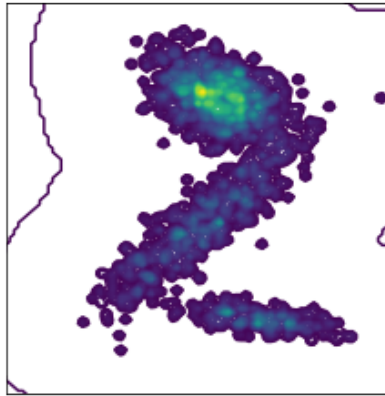
fig, axs = plt.subplots(len(h_list), 2, figsize=(8, 16),
    ↳ gridspec_kw=dict(left=0.1, right=0.9, bottom=0.1, top=0.96))
plt.subplots_adjust(top=0.8)
fig.suptitle('Parzen window density estimation', fontsize=18)

for i, h in enumerate(h_list):
    model_gk = ParzenWindow(X_train, h, 'gaussian')
    density_g = model_gk.predict(coordinates).reshape(x_grid.shape)

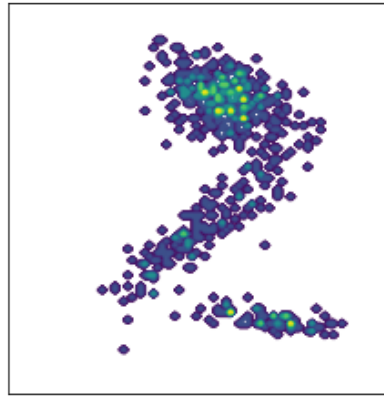
    model_wk = ParzenWindow(X_train, h, 'window')
    density_w = model_wk.predict(coordinates).reshape(x_grid.shape)

    ax1, ax2 = axs[i,0], axs[i,1]
    ax1.contour(x_grid, y_grid, density_g, levels=np.linspace(density_g.min(),
    ↳ density_g.max(), num=100))
    ax1.set_xlabel(f'Gaussian kernel, h={h}', fontsize=13)
    ax1.set(xticks=[], yticks=[])
    ax2.contour(x_grid, y_grid, density_w, levels=np.linspace(density_w.min(),
    ↳ density_w.max(), num=100))
    ax2.set_xlabel(f'Window kernel, h={h}', fontsize=13)
    ax2.set(xticks=[], yticks=[]);
plt.show()
```

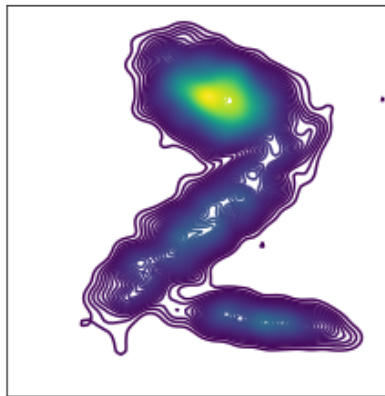
Parzen window density estimation



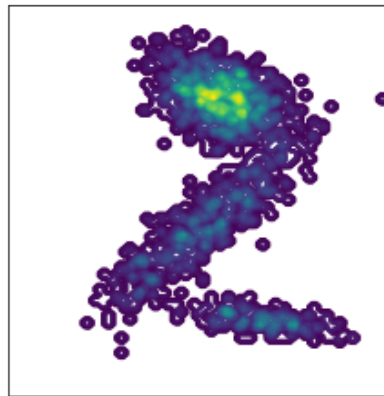
Gaussian kernel, $h=0.3$



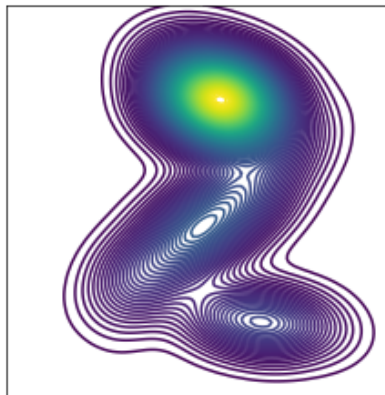
Window kernel, $h=0.3$



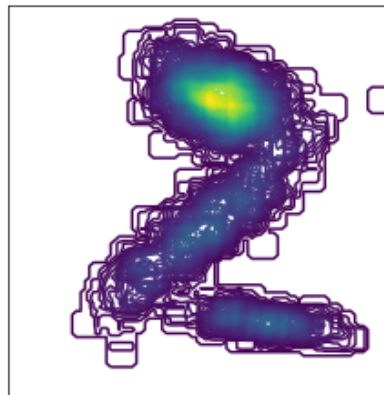
Gaussian kernel, $h=1$



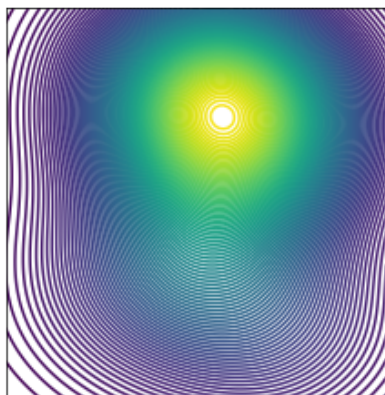
Window kernel, $h=1$



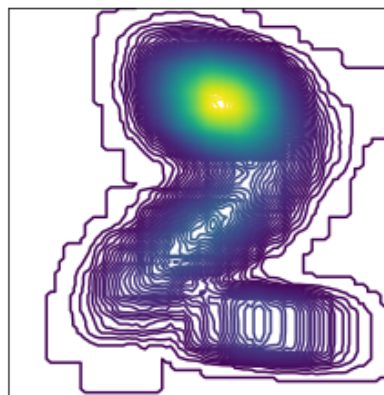
Gaussian kernel, $h=3$



Window kernel, $h=3$



Gaussian kernel, $h=9$



Window kernel, $h=9$

- As the window size increases, the density function becomes smoother and smoother.
- Small window sizes => density function changes over small changes in position
- Large window sizes => even if we add points in a small region, a larger surrounding will have its density function value changed because of the large window size. The contour plot is a graphical representation of the same.

3.1 GD GMM

- $\pi_k \in [0, 1]$ and $\sum \pi_k = 1$ then we declare a unconstrained variable ξ_k such that $\pi = \text{softmax}(\xi)$
- The covariance matrices must be positive semidefinite. so we declare unconstrained matrices M_k such that $\Sigma_k^{-1} = M_k M_k^T$

```
[13]: np.random.seed(42)
class GaussianMixtureModelGrad:
    def fit(self, X_train):
        self.si = np.zeros((Config.K,1)) # pi_k
        # pi is softmax(si)

        idx = np.random.choice(np.arange(Config.N), Config.K)
        self.mu = X_train[idx] * 1.0 # choose K samples randomly from dist
        # correct mu
        # self.mu = np.array([[3,3], [10, -10], [5, 20]]).astype(np.float32)
        self.precision = np.array([np.random.randn(Config.D, Config.D) for i in
        range(Config.K)]) # init K covs to identity
        self.eta = 2e-3/Config.N

        nll = []
        for i in range(Config.iterations*50):
            loss = self.gd_one_iteration(X_train)
            nll.append(loss)
        return nll

    # helpers
    def prob(self, X_nd, k): # (examples, k) -> [P(example1 | k),
    P(example2|k), ..]
        mu_1d = self.mu[k:k+1]
        # inv_cov_dd = self.precision[k] @ self.precision[k].T
        inv_cov_dd = np.linalg.inv(Actual.cov[k])
        z_nd = X_nd - mu_1d
        A_nd = z_nd @ inv_cov_dd
        z_n = (A_nd * z_nd).sum(axis=1)
        p_n = (2*np.pi) ** (- Config.D / 2) * np.linalg.det(inv_cov_dd) ** .5
        # np.exp(-.5 * z_n)
        return p_n
```

```

def getGamma(self, X_nd):
    g_kn = np.array([
        np.exp(self.si[k, 0]) * self.prob(X_nd, k)
        for k in range(Config.K)
    ]) # (k n)
    denom = g_kn.sum(axis=0, keepdims=True)
    denom[denom == 0.0] = 1
    g_kn /= denom # sum over clusters (k) must be one

    return g_kn, denom

def gd_one_iteration(self, X_nd):
    """Performs one iteration of the gradient descent for gaussian mixture_
    ↪model
    Args
        * X_nd is an np.array with n examples, each of d dimensions
    Function
        * updates Params.mu, Params.cov and Params.pi
    """
    g_kn, denom = self.getGamma(X_nd)

    # N_k is gamma summed over n
    N_k1 = g_kn.sum(axis=1, keepdims=True) # shape (k,1)

    # new mu is examples averaged over gamma
    for k in range(Config.K):
        z_nd = X_nd - self.mu[k:k+1]
        g_n1 = g_kn[k:k+1].T
        # inv_cov = self.precision[k] @ self.precision[k].T
        inv_cov = np.linalg.inv(Actual.cov[k])
        delta_mu = g_n1.T @ (inv_cov @ z_nd.T).T

        # delta_cov = .5 * (inv_cov @ ((g_n1 * z_nd).T @ z_nd) @ inv_cov -
        ↪N_k1[k] * inv_cov)
        # delta_precision = np.linalg.inv(self.precision[k].T) * N_k1[k] -
        ↪self.precision[k] @ ((g_n1 * z_nd).T @ z_nd)

        # self.precision[k] += self.eta * delta_precision
        self.mu[k:k+1] += self.eta * delta_mu

    delta_si = N_k1
    self.si += self.eta * delta_si

    # return log likelihood
    neg_log_likelihood = - np.log(denom).sum()
    return neg_log_likelihood

```

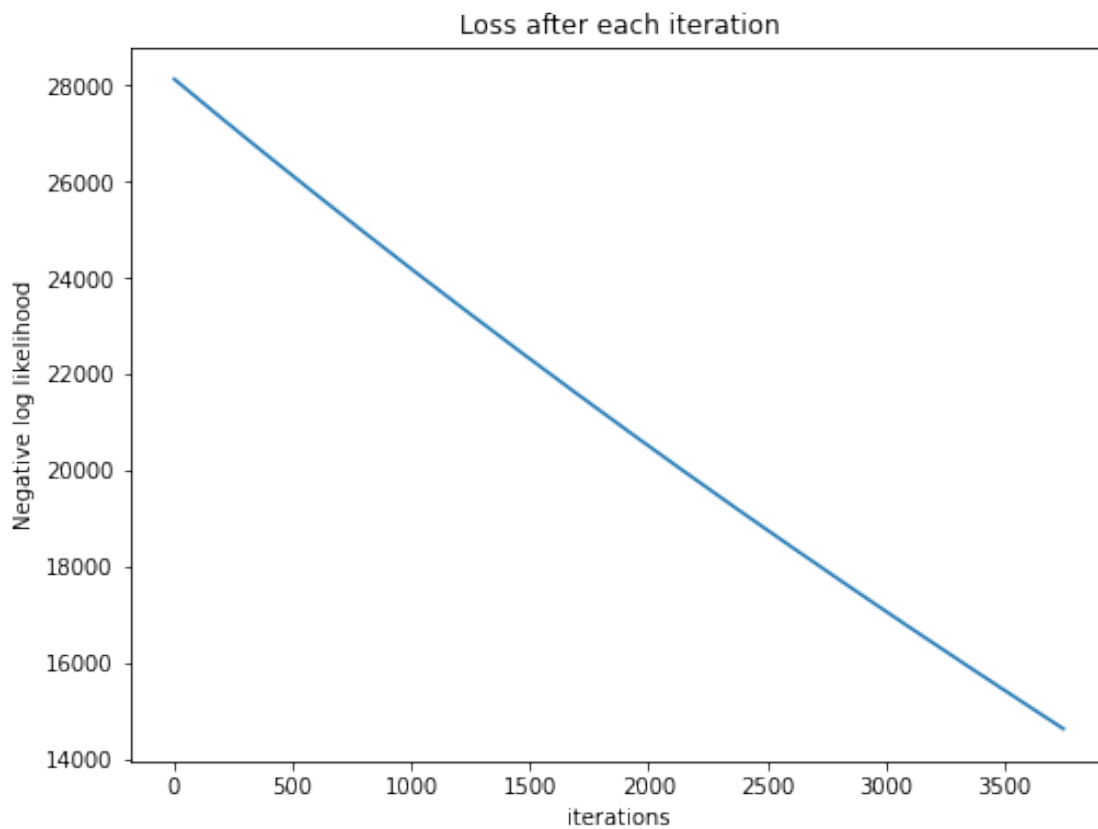
```

model = GaussianMixtureModelGrad()
nll = model.fit(X_train)
fig, ax = plt.subplots(figsize=(8, 6))
plt.plot(nll)
plt.xlabel('iterations')
plt.ylabel('Negative log likelihood')
plt.title('Loss after each iteration')
plt.show()

print("Model means:")
print(model.mu)
print()
print("Actual means:")
print(Actual.mu)

print(model.si)

```



```

Model means:
[[ 9.56550783 17.95945598]
 [ 4.56785114 15.78477678]]

```

```
[ 6.97719713 18.4234213 ]]
```

Actual means:


```
[[ 3  3]
 [10 -10]
 [ 5 20]]
[[2.89094286]
 [0.40089697]
 [4.20816017]]
```

```
[14]: fig, ax = plt.subplots(figsize=(10, 10))
      parts = np.cumsum([0] + Actual.sizes)

      x_grid = np.arange(-20, 25, .02)
      y_grid = np.arange(-20, 27, .02)
      x_grid, y_grid = np.meshgrid(x_grid, y_grid)
      coordinates = np.array([x_grid.ravel(), y_grid.ravel()]).T
      colors = ['green', 'orange', 'blue']

      for k in range(Config.K):
          x = X_train[parts[k]:parts[k+1]]
          ax.scatter(x[:,0], x[:,1], label=str(k+1))
          cov = np.linalg.inv(model.precision[k] @ model.precision[k].T)
          z_grid = scipy.stats.multivariate_normal(model.mu[k], cov).pdf(coordinates).
          ↪reshape(x_grid.shape)
          ax.contour(x_grid, y_grid, z_grid, colors=colors[k])

      ax.scatter(model.mu[:,0], model.mu[:,1], marker='*', label='centers')
      ax.set_title('After fitting GMM')
      plt.legend();
```



```
[15]: nll[-10:]
```

```
[15]: [14645.892619211156,  
      14642.69083918832,  
      14639.489220479289,  
      14636.287763098833,  
      14633.086467061868,  
      14629.885332383456,  
      14626.684359078783,  
      14623.483547163189,  
      14620.282896652141,  
      14617.082407561255]
```



```
[16]: Actual.cov
```

```
[16]: [array([[34., 30.],  
            [30., 34.])),  
      array([[15.54769466, -2.56515107],  
            [-2.56515107,  1.45230534]]),  
      array([[13.5      , -4.33012702],  
            [-4.33012702,  8.5      ]])]
```